

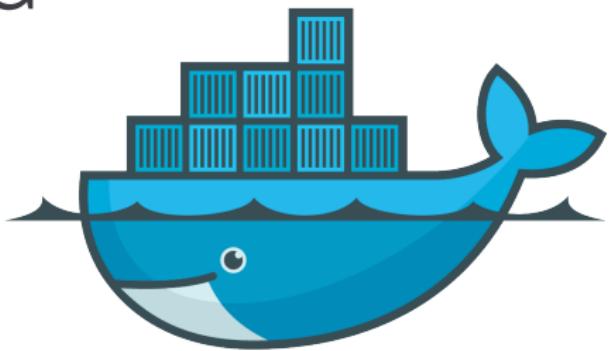
HANDS-ON WITH DOCKER FOR DEEP LEARNING

Taming the whale

01/12/2016

Marc-André Gardner

Université Laval



docker

Why Docker?

Let's play with deep learning!

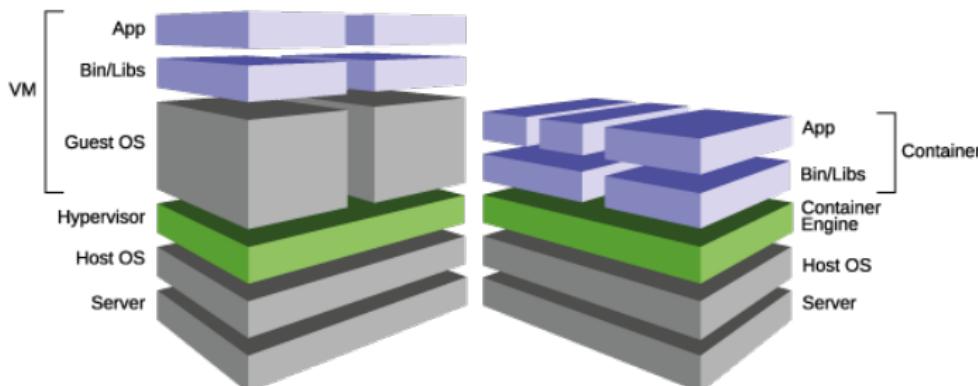
- Ok, let's try Tensorflow!
 - But not r0.12! I need r0.10!
 - Ok but r0.10 is incompatible with cuDNN R5, I need cuDNN r3!
 - Hum right but then I cannot use CUDA 8.0
 - But my brand new GeForce 1080 needs that driver version!
- You know what, let's start with this Caffe model instead.
 - Fine, but now I need Boost 1.55.
 - And protobuf, hdf5lib and Imdb
 - And glog
 - What is glog?
 - Ok everything is installed, can I *finally* run something?
 - Not so fast, the code you want to use is built on an old, never released, caffe dev branch
 - (3 hours browsing the Internet, struggling to retrieve this branch)
 - Surprise! This branch does not support cuDNN R5 either!
- Hmm what about using Torch instead? Let's install it
 - Ok, first I must install Lua
 - Damn, its OpenBLAS is conflicting with the one I installed for numpy earlier
 - Wait, I have built the Torch executables on another similar workstation.
 - HAHA! Your graphics cards do not share the exact same streaming multiprocessors version (whatever that means)! Now everything crash...

After a few hours like that...



What is Docker?

- Docker allows the deployment of *software containers*
- Each container is *standalone*, and has its own:
 - Filesystem
 - Libraries
 - Configuration files
- Each container is also *isolated* from the others and the host
- The only shared part is the kernel
- Can be seen as a high-level virtual machine



What is useful with Docker?

- Docker allow us to see the whole system as a code repository:
 - You can *commit* any change you made to the system
 - If there is an error, you can *revert* them
 - Also, this error will not affect the host system
 - If you want to run the *exact same* system configuration after 6 months without using it, you can
 - You can *bundle* your system, and then:
 - You can *backup it*
 - You can *transfer it* to another computer
 - You can *share it* with others
 - You can retrieve actual systems from online repositories
 - It is *your system*: in the Docker container, you are the indisputed master!
 - You can run more than one system at the same time
- Since the only shared part is the Linux kernel, you may run different distributions (Ubuntu, Fedora, Centos, etc.) on the same machine at the same time!

What Docker is not?

- Docker does not ease the libraries installation process
 - Except for the fact that you may always take one step back if you made a mistake
- Docker is not a *data* backup solution: keep your datasets and your results safe!

Images vs Containers

- An image is like a music score
- A container is what is *executing* this score (namely, an orchestra)
- We can copy the score and give it to another orchestra
- In the Docker world, we may also pause the orchestra, record it (so to make a new score) or terminate it at any moment.
 - Do not try this in real life



Images vs Containers

- A Docker **image** contains a *filesystem state*
- You can (and should) **tag** an image
 - Multiple tags can be added
- A Docker **container** contains a *running system state*
- A **Dockerfile** is a text file containing the instructions needed to *build* a Docker



Docker images

- An image is identified by its image ID, like 4201baf875aa
- One can assign a repository and a **tag** to an image:

```
$ docker tag 4201baf875aa ubuntu-cudnn:caffe
```

- The name must only contains lowercase letters, digits, _ and -
- Usually, an image name is made of two parts, separated by a colon, like “repository:tag”

- An image can be loaded from a local file:

```
$ docker load -i /path/to/my/imagefile.tar
```

- Conversely, it can be saved to a local file:

```
$ docker save -o /path/to/my/imagefile.tar imageid
```

- An image can also be pulled from a remote server:

```
$ docker pull repository:tag
```

- You can view the currently loaded images plus some info about them:

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
ubuntu-cudnn	caffe	4201baf875aa	6 weeks	4.874 GB

Docker containers

- Once we have an image, we can execute it in a **container**:

```
$ docker run IMAGE COMMAND
```

- Usually, we just want to run a terminal in our image so we can start our task ourselves:

```
$ docker run IMAGE /bin/bash
```

- Hey! Nothing happened, the command just returns!

- You have to ask Docker to behave like a TTY, or the shell will think that it is parsing a file and exit after reaching EOF

```
$ docker run -it IMAGE /bin/bash
```

- You may list the containers using:

```
$ docker ps -a
```

CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	NAMES
fc655d9d3f6e	ubuntu	<code>/bin/bash"</code>	3 seconds ago	Up 2 sec	dreamy_dijkstra
37a636090f3f	ubuntu	<code>/bin/bash"</code>	3 minutes ago	Exited (0)	awesome_poitras

- A container remains even after being closed!

- To avoid that, use the `-rm` option:

```
$ docker run -it --rm IMAGE /bin/bash
```

More fun with containers

- A container can be **exported**:

```
$ docker export -o /path/to/my/file.tar CONTAINER
```

- Where CONTAINER is either the container ID or the name of the container
- **Not to be confounded with** saving an image!

- The produced tarball may then be imported

```
$ docker import /path/to/my/file.tar
```

- You can pause an entire container at any moment:

```
$ docker pause CONTAINER
```

```
$ docker unpause CONTAINER
```

- The changes made in a container are *temporary*; when you exit, the underlying image stays unaffected

- If an orchestra ruins a song, it does not destroy the song itself!

- But if you *want* your modifications to be permanent:

```
$ docker commit -m MESSAGE CONTAINERID [IMAGEID]
```

- This will create a new image

Inside a container

- Starting a container with `/bin/bash` as executable target gives you a prompt:

```
$ docker run -it --rm ubuntu /bin/bash
root@87132209cb86:/#
```

- Wait, I have root access?
- Yes, but only in the context of this very container; you cannot access to anything outside it!

- You cannot even see the other processes running on the host

```
root@87132209cb86:/# ps aux
USER  PID %CPU %MEM VSZ    RSS    STAT START   COMMAND
root   1  0.0  0.0  18248   3220   Ss  21:23  /bin/bash
root  15  0.0  0.0  34424   2824   R+  21:32  ps aux
```

- Now you may install everything you need, as a super-user!
- And if you break something?
 - Just exit and restart the container
 - As long as you did not *commit* anything, the changes you made are temporary

Peaking outside: using Docker volumes

- The isolation of each container is great, but what if I want to access outside files?
- One solution could be to put and commit them in the image itself
 - Not very interesting for big datasets (use a lot of space)
 - Not interesting either for code, as you have to make an update from the code repo each time you are starting a new container...
- Another better way: Docker **volumes**
- Use the option `-v path:mountpoint:permissions`
 - path is the path on your host system you want to share
 - mountpoint is the path of the mountpoint on the container. If it does not exist in the image, Docker will create it.
 - permissions is either ro (read-only) or rw (read and write)
- \$ docker run -it --rm -v \$HOME/mycode:/mnt/code:ro ubuntu /bin/bash
root@87132209cb86:/# ls /mnt/code
createdataset.py deepsn_batch.py observeLoss.lua train.py
- You can mount multiple volumes (provide many -v arguments)

Using a GPU

- We would like to use GPUs inside Docker containers.
- In itself, it is not hard :

```
$ docker run --devices=/dev/nvidia*:/dev/nvidia* ...
```
- But then you must install everything in your image
- Since the nVidia driver and CUDA are specific to each GPU architecture, we are in for some fun...
- Fortunately, nVidia has our back covered with nvidia-docker:

```
$ nvidia-docker run -it --rm IMAGE /bin/bash
```
- No need to use nvidia-docker anywhere else (though it works)

Other useful tools

- `docker rm` deletes a *container* (if it is running, you must explicitly use the `-force` flag to kill it)
- `docker rmi` deletes an *image*
- You can **attach** yourself to any running container, just use:
`$ docker attach CONTAINERID`
- To **detach** yourself without aborting the execution of the container, press `Ctrl-p` and then `Ctrl-q`
- If you *create* files on a volume, you will notice that they are all owned by root
 - It is not a huge problem since they are still readable and you may delete them from a container, but it is annoying
 - To avoid that, you may pass a userid to `docker-run`:
`$ docker run -u `id -u $USER` -it --rm ubuntu /bin/bash`
I have no name!@f7cb7055a581:/\$
 - You get a weird prompt, but now your user on the host owns the file created by the container
- `docker stats` and `docker top` are useful to remotely check the status and resource usage of all containers

Cheat sheet

